ViMed

Scheduled micro-consultations between healthcare professionals & pharma reps

Project Compass

Product & market brief for the ViMed web prototype — a feedback-driven proof of concept for the Saudi Arabian market.

Prototype · Web (Next.js) | Market: Saudi Arabia | Bilingual · Arabic-first / RTL | Purpose: testing & surveys

Prepared as a business & product compass | Companion to the technical Build Plan & Architecture documents

CONTENTS

What's inside

> **HOW TO READ THIS**
>
> This Compass is the **business & product** view — written for founders, stakeholders, and the people you'll survey. The accompanying **Build Plan** tells Claude Code how to build it step by step, and the **Architecture & Decisions** document explains the engineering reasoning. All three describe the same product.

The idea in one page

ViMed is a scheduling platform that connects **Healthcare Professionals (HCPs)** — physicians, pharmacists, and purchasers — with **pharmaceutical sales representatives** through short, pre-booked **120-second video calls**. Instead of reps competing for unplanned hallway minutes, HCPs publish the slots they're willing to give, and reps book them respectfully and on the record.

The wedge

Replace the chaotic, gate-kept "drop-in" sales visit with consented, time-boxed, scheduled micro-meetings.

Why short calls

120 seconds forces signal over noise — one product, one message — and respects clinical time.

Who it's for

The Saudi healthcare ecosystem: MOH & private hospitals, retail & hospital pharmacies, and the reps who serve them.

Why now

This prototype exists to **test demand and gather structured feedback** before committing to a full build.

What this prototype is — and isn't

This is a **web prototype built for learning**, not a production launch. It demonstrates the complete experience end-to-end so real HCPs and reps can use it, react to it, and tell us what works. Some pieces are deliberately simulated (notably the video call) so we can validate the *concept and the flow* without the cost and complexity of live telephony. Every simulated piece is built behind a clean seam so it can be swapped for the real thing later without rework.

> **THE SINGLE GOAL OF THIS BUILD**
> Put a believable, bilingual, KSA-flavoured ViMed in front of real users and **collect feedback** — via in-app ratings, an admin analytics view, and linked external surveys — to decide whether the business idea deserves a full investment.

The core loop, in one breath

1. **HCP sets availability** — a doctor or pharmacist publishes the weekly days & hours they'll accept calls, sliced into 15-minute slots.
2. **Rep finds & books** — a rep filters the directory (city, sector, centre, specialty), opens a profile, and books an open slot.
3. **They meet for 120 seconds** — at the appointed time, a simulated in-browser call runs a 2-minute countdown.

4. **Both rate the visit** — a 5-star + comment exchange feeds reputation, loyalty points, and our feedback data.

Why ViMed exists

The pain on both sides

For pharma reps

Access to busy clinicians is the hardest part of the job. Reps spend hours travelling and waiting for a few unscheduled minutes that may never come. Hospital access is gated, time is unpredictable, and there's no reliable record of who was seen, when, or what was discussed.

For HCPs

Physicians and pharmacists are interrupted at unpredictable moments by reps competing for attention. They have no clean way to say "I'll give you two minutes on Tuesday at 9:15" and have it respected — and no easy way to keep up with genuinely useful product and treatment updates.

The opportunity

ViMed turns an adversarial, low-trust interaction into a **consented, scheduled, measured** one. The HCP stays in control of their time; the rep gets guaranteed, qualified attention; and every interaction generates structured data — ratings, frequency, topics — that neither side has today.

> **THE INSIGHT**
> The bottleneck isn't *willingness* to meet — it's the absence of a trusted, low-friction way to *schedule and time-box* the meeting. Solve scheduling and consent, and the rest follows.

Adjacent value we layer on

- **A medical library** — curated updates plus a searchable treatment-protocol reference (FDA/NHS-style guidelines), giving HCPs a reason to return between calls.
- **Reputation** — mutual 5-star ratings build accountability on both sides.
- **Loyalty & gamification** — points, redeemable offers, and a leaderboard that nudge healthy engagement.

Built for Saudi Arabia

The prototype is scoped to the Saudi healthcare market across three anchor cities — **Riyadh, Jeddah, and Dammam** — spanning both the **Ministry of Health (MOH)** and the **private sector**.

The directory we're modelling

Audience segment	Count (mock)	Detail
Physicians	100	10 specialties × 10 — Dermatology, Orthopedics, Cardiology, Dentistry, Pulmonary, Internal Medicine, Pediatrics, Endocrinology, General Surgery, OB/GYN
Pharmacists	30	Retail (AlNahdi, AlDawa, UPC) & hospital (MOH + private)
Purchasers	10	MOH hospitals & private hospital groups
Sales reps	seeded set	Tied to companies/agents, territories, specialties & product/SKU lists

Centres include MOH sites (King Fahd Hospital, King Faisal Hospital, Prince Sultan Medical City, Jeddah Central Hospital) and private sites (Dr. Sulaiman Al-Habib, Al Hammadi, Saudi German, Tadawi Clinics, Nahdi, Al Dawa). Doctors, pharmacists and purchasers are randomly assigned across centres and cities.

Saudi-specific design requirements

Arabic-first & RTL

Arabic is the **default language**, with a one-tap switch to English. The entire interface mirrors correctly in right-to-left: layout, navigation, icons, form fields, and the calendar all flip — not just translated text.

Bilingual content model

Personal names are presented in **Arabic**; centre and city names stay in **English** as they're commonly written. Every label exists in both languages from day one.

Mobile-first reality

The Saudi audience is heavily smartphone-led, so the web app is **designed mobile-first** and scales up gracefully to tablet and desktop.

Local structure

Sector (MOH vs Private), territory, and centre taxonomy mirror how the Saudi market actually segments — so filters feel native to a local rep or clinician.

LOCALISATION IS A FEATURE, NOT A TRANSLATION PASS

For this audience, credibility depends on the app *feeling* Saudi from first launch — Arabic default, correct RTL behaviour, familiar institutions, and local naming. It's treated as core product, not an afterthought.

Who uses ViMed

The prototype supports three roles. The two market-facing roles drive the core loop; Admin exists to observe it and gather feedback.

① Healthcare Professional (HCP)

Physician · Retail pharmacist · Hospital pharmacist · Purchaser

Goals

Control when and how often reps reach me; stay current on relevant products & protocols; be recognised for engaging.

Pains

Unplanned interruptions; no record of interactions; hard to filter signal from noise.

What's unique: only HCPs publish availability. They define duty days, set start/end hours per day, and the system generates bookable 15-minute slots.

② Pharmaceutical Sales Representative

Goals

Reliable, qualified access to the right HCPs; efficient territory coverage; a good reputation.

Pains

Wasted travel & waiting; gate-keeping; no proof of engagement.

What's unique: reps search & filter the directory, book slots, reschedule/cancel, start the call, and rate the HCP afterwards.

③ Administrator

Goals

Observe how the prototype is used; read aggregated ratings & engagement; route testers to surveys.

Why it exists

This role is the **feedback engine** of the prototype — it turns usage into learning.

> ROLE ASYMMETRY TO REMEMBER
>
> **HCPs supply time; reps consume it.** Availability authoring is HCP-only; booking against open slots is rep-only. Admin neither supplies nor consumes — it measures.

Everything the prototype does

Extracted and re-platformed from the original brief and wireframes. Grouped by area; each line notes whether it's primarily a frontend or backend concern (most are both).

Accounts & access

- Separate sign-up / sign-in journeys for **HCP** and **Sales Rep** FE BE
- Email + password authentication with secure sessions (JWT); seeded demo accounts for testers BE
- Role-based access — every screen and action gated by role BE
- Profile capture: name, gender, DOB, occupation, degree, specialty/sub-specialty, centre, company/products (rep) FE
- Terms & conditions consent FE

Directory & discovery (rep)

- Filters: territory, sector (MOH/Private), centre, specialty FE BE
- Search by doctor name or centre name BE
- Profile view with available time slots FE

Availability & scheduling (HCP)

- Weekly availability authoring — choose duty days, set hours per day FE BE
- Auto-generated 15-minute slots from chosen hours BE
- "Fixed weekly" repeat option BE
- Calendar month/day views & navigation FE

Visits & the call

- Book a visit against an open slot BE
- Upcoming & History sub-views FE
- Reschedule & cancel (rep) BE
- Start a **simulated 120-second video call** (local camera + countdown) FE
- Post-call mutual 5-star rating + comment BE
- Real-time updates & reminders via WebSockets BE

Library

- Medical Updates — news, articles, launches
- Treatment Library — protocol search with specialty/diagnosis filters

Loyalty

- Points tracker
- Redeemable offers
- Leaderboard (gamification)

Admin & feedback

- Analytics dashboard (visits, ratings, active users by role/city/specialty)
- Leaderboard & loyalty overview

Platform-wide

- Arabic ⇄ English with full RTL
- Mobile-first responsive app shell
- "Soft Care" themed design system

- External survey links + click tracking

Every screen, mapped

The app shell uses a primary navigation of five areas (mirroring the original mobile tabs), plus authenticated sub-routes. Availability authoring is HCP-only; booking is rep-only; Admin lives in its own area.

ViMed

- Public
  - Landing / about the test
  - Sign in
  - Sign up → HCP path · Sales Rep path
  - Language switch (AR ⇄ EN)
- App shell (authenticated)
  - Visits
    - Upcoming
    - History
    - Visit detail → (rep) reschedule · cancel · start call
  - New Visit (rep)
    - Directory + filters (territory / sector / centre / specialty)
    - Search results
    - HCP profile → slot picker → confirm booking
  - Availability (HCP)
    - Weekly day & hours editor
    - Fixed-weekly toggle
    - Generated slots · month/day calendar
  - Library
    - Medical Updates
    - Treatment Library (search + filters)
  - Loyalty
    - Points tracker
    - Offers
    - Leaderboard
  - Profile
    - View / edit info
    - Ratings received
    - Settings (language) · sign out
  - Call
    - Simulated call room (120s)

- Post-call rating

Admin area

- Dashboard (KPIs)
- Visits & ratings analytics
- Users by role / city / specialty
- Survey links & click tracking

Role-based journeys

HCP — from sign-up to being bookable

1. **Sign up as HCP** — choose occupation & degree, enter specialty, select city & centre from dropdowns, accept T&Cs.
2. **Publish availability** — pick duty days, set start/end hours per day; optionally mark "fixed weekly" to repeat.
3. **Slots appear** — the system generates 15-minute slots; the HCP confirms which to offer.
4. **Get booked** — a rep books a slot; the HCP sees it under Upcoming and receives a real-time notification.
5. **Take the call & rate** — join the 120-second simulated call, then rate the rep. Earn loyalty points.

Sales Rep — from search to rated visit

1. **Sign up as Rep** — enter company/agent, specialties covered, products & SKUs, accept T&Cs.
2. **Find an HCP** — open New Visit, filter by territory/sector/centre/specialty or search by name.
3. **Book a slot** — open a profile, pick an open 15-minute slot, confirm the visit.
4. **Manage it** — reschedule or cancel from Upcoming if plans change (real-time sync to the HCP).
5. **Start the call & rate** — at the slot time, start the simulated call; afterwards, rate the HCP.

Admin — from observation to insight

1. **Sign in to Admin** — access the dedicated analytics area.
2. **Read the dashboard** — visits booked/completed, average ratings, active users by role/city/specialty, loyalty leaderboard.
3. **Route to surveys** — surface external survey links to testers and watch click-through.
4. **Decide** — combine in-app signals + survey responses to judge product–market fit.

"Soft Care"

A calm, approachable, patient-friendly medical aesthetic — softer than clinical blue, warmer than corporate green. It signals trust and care without feeling sterile, and it reads beautifully in both Arabic and English.

Palette

- **Sage — #6FA287** · primary: actions, navigation, brand
- Sage Deep — #557C68 · headings, emphasis, hover
- Sage Light — #A7C6B5 · borders, dividers, subtle fills
- **Aqua — #9AD0DD** · accent: highlights, secondary actions
- Aqua Deep — #6FB3C4 · links, accent emphasis
- Warm White — #FBFAF6 · app background
- Ink — #2B3A33 · primary text

Voice & feel

Tone

Reassuring, professional, uncluttered. Plenty of whitespace, gentle rounded corners, soft shadows.

Typography

An Arabic-first type system (e.g. Cairo / Tajawal for Arabic, Inter for Latin) so both scripts feel intentional, not bolted on.

Imagery

Light, human, healthcare-positive — never alarming or overly clinical.

Accessibility

Calm contrast that still meets readability standards; large, finger-friendly tap targets for mobile.

Real vs mocked

Smart prototyping means building the parts that need to feel real, and simulating the parts that are expensive to validate — always behind clean seams so the simulated parts can be upgraded later without a rewrite.

Area	In this prototype	Later (production)
Authentication	**Real** — email + password, JWT, seeded demo logins	Add phone/OTP, email verification, password reset
Directory	**Realistic mock data** — 100 physicians, 30 pharmacists, 10 purchasers, reps; KSA centres & cities	Real onboarding & verification of HCPs/reps
Scheduling	**Real** — availability authoring, 15-min slot generation, booking	Calendar sync, conflict handling at scale
Video call	**Simulated** — local camera + 120s countdown → rating	Real WebRTC (e.g. LiveKit / Daily / Agora)
Real-time	**Real** — WebSocket updates & notifications	Scale-out, push notifications
Maps / territory	**Dropdowns** — city → centre selection (no map)	Map-based territory drawing & geolocation
Library	**Mock content** — sample updates & protocols	Real curated/licensed guideline feeds
Loyalty	**Functional demo** — points, offers, leaderboard	Real rewards partners & redemption
Surveys	**External links** + click tracking	Embedded research instrumentation

DESIGN PRINCIPLE

Every mocked capability is isolated behind an interface (a "provider" or "adapter"). Swapping the simulated call for real WebRTC, or mock content for a live feed, should be a contained change — not a refactor.

What the test should learn

The prototype's whole reason for existing is to produce evidence. Three feedback channels feed the decision.

① In-app ratings

The mutual 5-star + comment after every call — a continuous read on perceived value and behaviour quality.

② Admin analytics

Usage truth: how many visits get booked vs completed, average ratings, who's active by role/city/specialty, leaderboard movement, survey click-throughs.

③ External surveys

Deeper qualitative & quantitative feedback via linked forms, surfaced in-app and tracked.

Questions the data should answer

- **Demand:** Do HCPs actually publish availability? Do reps actually book?
- **Value:** Are 120-second calls perceived as worthwhile by both sides?
- **Friction:** Where do users drop off — sign-up, availability, booking, the call?
- **Stickiness:** Does the Library / Loyalty layer bring people back?
- **Fit:** Which segments (specialty, sector, city) respond best?

SUGGESTED KPIS TO WATCH

Availability publish rate · booking conversion · visit completion rate · average mutual rating · repeat usage · survey completion · qualitative themes from comments.

Prototype → pilot → production

Now — Prototype (this build)

- Bilingual web app, real auth & scheduling, simulated call, mock directory & content.
- Admin dashboard + survey links for feedback collection.
- Goal: validate demand and the core loop with real Saudi HCPs & reps.

Next — Pilot (if signals are positive)

- Swap simulated call for real WebRTC; add phone/OTP & verification.
- Onboard a small set of real HCPs & reps in one city/sector.
- Replace mock content with a curated guideline feed.

Later — Production

- Cloud deployment & scale-out, push notifications, calendar sync.
- Real loyalty/reward partners; map-based territory tools.
- Compliance, analytics, and commercial model build-out.

BOTTOM LINE
Build the smallest believable ViMed, put it in front of real users, and let the feedback — not the assumptions — decide what comes next. The architecture is deliberately built so that "what comes next" is an upgrade, never a restart.

ViMed

Project Compass · companion to the Build Plan & Architecture & Decisions documents